

# NoSQL

(Or “Not Only SQL”?)

Databases that aren't  
based on tabular models  
and SQL

Many different models:  
Document, Key Value,  
Column, Graph

Advantages

Performance

# Scaling

Data model

OOP



Agile

Comes at a cost

# CAP Theorem

- Impossible for a distributed system to simultaneously provide all three:
  - **C**onsistency
  - **A**vailability
  - **P**artition Tolerance

NoSQL usually prefers  
the last 2 and have  
**eventual consistency**

NoSQL models

# Column based

- Column groups almost look like tables
- But queries are very simplified (e.g. only query by primary key)
- Allows for massive distribution and scale
- Examples:
  - Cassandra
  - HBase

# Key-Value stores

- A lot like huge hash maps
- Bunch of keys, each key has its own values (numbers, strings, lists, etc.)
  - key: “user:1000”, value: {name: ‘Moshe’}
  - key: “user:1001”, value: {name: ‘Michael’}
  - key: “project:2”, value: {owner: 1000}
- Examples:
  - Redis
  - Riak

# Graph databases

- Let's you have nodes, links and properties
- For example: Facebook (person is a node, friendships are links)
- Examples:
  - neo4j



# Document based

- Documents are the building blocks - basically objects
- There's no schema
- Comfortable for ORM
- Very nice for REST
- Examples:
  - MongoDB
  - CouchBase

MongoDB

# MongoDB

- One of the most popular NoSQL databases
- JS interface
- Great replication
- Great sharding
- Easy to start working with
- Strong queries

A database contains  
many collections

e.g. Users, Projects, Products

Each collection contains documents which are a lot like JSON objects

Documents can contain  
strings, numbers, lists,  
nested objects, etc.

No joins

Operations



# Inserting

```
db.collection('Todos').insertOne(todo, function(err, result) {  
  if (err) {  
    // handle error  
    return;  
  }  
  
  var savedTodo = result.ops[0];  
});
```

# ObjectID

- Automatically created for inserted objects under the key `_id`
- Look like this:  
`ObjectID('549d6cbecac7b55bc0a9c7fb')`

# Fetching single document

```
db.collection('Todos').findOne({_id: ObjectId(todoId)}, function(err, todo) {  
  if (err) {  
    // handle error  
    return;  
  }  
  
  if (todo === null) {  
    // no todo with specified id found  
  } else {  
    // use our todo  
  }  
});
```

# Fetching multiple documents

```
db.collection('Todos').find({done: true}).toArray(function (err, completedTodos) {  
  if (err) {  
    // handle error  
    return;  
  }  
  
  // Use our completedTodos  
});
```

# Updating a document

```
db.collection('Todos').updateOne({_id: ObjectId(todoId),
  {$set: {done: true}}, function(err, result) {
    if (err) {
      // handle error
    }

    // our update has finished
  });
```

# Queries

AND

{done: false, user: 'Aviv'}

OR

{\$or:

[{done: false}, {user: 'Aviv'}]}



```
{  
  done: false,  
  $or: [{user: 'Aviv'}, {user: 'Elad'}]  
}
```

{users: {\$gt: 3}}

{users: {\$lt: 10}}

```
{user: {sin: ['Aviv',  
             'Elad']}}
```

```
{  
  users: {  
    $elemMatch: {  
      $in: ['Aviv', 'Elad']  
    }  
  }  
}
```

{counts: {\$all: {\$gt: 5}}}

# Lots more

- \$size
- \$regex
- \$where
- \$exists
- And nicer functions (findAndModify)

# Bringing promises to Node

Q library



Q.when(value);

```
Q.all([promise1,  
      promise2]);
```

```
Q.ninvoke(db.collection('Todos').find({}), 'toArray').then(  
  function(todo) {  
    // do something with todo  
  }  
)  
.fail(  
  function(err) {  
    // handle error  
  }  
);
```